

APPM 4650 — HOMEWORK # 2 Solutions

1. (a) Show that $(1+x)^n = 1 + nx + o(x)$ as $x \rightarrow 0$.
- (b) Show that $x \sin \sqrt{x} = O(x^{3/2})$ as $x \rightarrow 0$.
- (c) Show that $e^{-t} = o(\frac{1}{t^2})$ as $t \rightarrow \infty$.
- (d) Show that $\int_0^\varepsilon e^{-x^2} dx = O(\varepsilon)$ as $\varepsilon \rightarrow 0$.

Soln:

- (a) In this problem $f(x) = (1+x)^n - (1+nx)$ and $g(x) = x$. Then we need to look at the $\lim_{x \rightarrow 0}$ of $\frac{f(x)}{g(x)}$. If it goes to 0, we have little o.

$$\begin{aligned} \lim_{x \rightarrow 0} \frac{(1+x)^n - (1+nx)}{x} &= \lim_{x \rightarrow 0} \frac{n(1+x)^{n-1} - n}{1} \quad \text{by L'Hopital's rule} \\ &= n - n = 0 \end{aligned}$$

- (b) In this problem $f(x) = x \sin(\sqrt{x})$ and $g(x) = x^{3/2}$. We need to find a positive constant M such that $\left| \frac{f(x)}{g(x)} \right| \leq M$.

$$\begin{aligned} \left| \frac{f(x)}{g(x)} \right| &= \left| \frac{x \sin \sqrt{x}}{x^{3/2}} \right| \\ &= \left| \frac{\sin \sqrt{x}}{x^{1/2}} \right| = \left| \frac{\sin \sqrt{x} - 0}{x^{1/2} - 0} \right| \\ &= \left\| \frac{\sin u - 0}{u - 0} \right\| \quad \text{where } u = \sqrt{x} \text{ (Note: when } x \text{ is near } 0, u \text{ is near } 0) \\ &= \|\cos \eta\| \text{ for some } \eta \in (0, u) \text{ by the Mean Value Theorem} \\ &\leq 1 = M \end{aligned}$$

- (c) In this problem $f(t) = e^{-t}$ and $g(t) = \frac{1}{t^2}$. Then we need to look at the $\lim_{t \rightarrow \infty}$ of $\frac{f(t)}{g(t)}$. If it goes to 0, we have little o.

$$\begin{aligned} \lim_{t \rightarrow \infty} \frac{e^{-t}}{\frac{1}{t^2}} &= \lim_{t \rightarrow \infty} \frac{t^2}{e^t} \text{ (simplification)} \\ &= \lim_{t \rightarrow \infty} \frac{2t}{e^t} \text{ by L'Hopital's rule} \\ &= \lim_{t \rightarrow \infty} \frac{2}{e^t} \text{ by L'Hopital's rule} \\ &= 0 \end{aligned}$$

- (d) In this problem $f(\varepsilon) = \int_0^\varepsilon e^{-x^2} dx$ and $g(\varepsilon) = \varepsilon$. We need to find a positive constant M such that $\left| \frac{f(\varepsilon)}{g(\varepsilon)} \right| \leq M$.

$$\begin{aligned}
\left| \frac{f(x)}{g(x)} \right| &= \left| \frac{\int_0^\varepsilon e^{-x^2} dx}{\varepsilon} \right| \\
&= \left| \frac{\int_0^\varepsilon e^{-x^2} dx - 0}{\varepsilon - 0} \right| \\
&= \|e^{-\eta^2}\| \text{ by the Mean Value Theorem } \exists \eta \in (0, \varepsilon) \text{ such that this is true} \\
&\leq 1
\end{aligned}$$

2. Consider solving $\mathbf{Ax} = \mathbf{b}$ where $\mathbf{A} = \frac{1}{2} \begin{bmatrix} 1 & 1 \\ 1 + 10^{-10} & 1 - 10^{-10} \end{bmatrix}$ and $\mathbf{b} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$. The exact solution is $\mathbf{x} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$ and the inverse of \mathbf{A} is $2 \begin{bmatrix} 1 - 10^{10} & 10^{10} \\ 1 + 10^{10} & -10^{10} \end{bmatrix}$. In this problem we will investigate a perturbation in \mathbf{b} of $\begin{bmatrix} \Delta b_1 \\ \Delta b_2 \end{bmatrix}$ and the numerical effects of the condition number.

- Find an exact formula for the change in the solution between the exact problem and the perturbed problem Δx .
- What is the condition number of \mathbf{A} ?
- Let Δb_1 and Δb_2 be of magnitude 10^{-5} . What is the relative error in the solution? How does the number of accurate digits relate to the condition number of \mathbf{A} and the perturbation in the right hand side?

Soln:

(a)

$$\Delta x = \mathbf{A}^{-1} \delta \mathbf{b} = \begin{bmatrix} \Delta b_1 - 10^{10}(\Delta b_1 - \Delta b_2) \\ \Delta b_1 + 10^{10}(\Delta b_1 - \Delta b_2) \end{bmatrix}$$

(b) $\kappa(\mathbf{A}) = 2 \times 10^{10}$

(c) I let $\Delta b_1 = 10^{-5}$ and $\Delta b_2 = 2 \times 10^{-5}$. Then the solution to the perturbed problem is $\begin{bmatrix} 10^4 \\ -10^4 \end{bmatrix}$. The relative error in the solution is 10^4 .

With a condition number of $O(10^{10})$ and relative perturbation of 10^{-5} we expect the relative error to be less than or equal to $O(10^5)$ which is what we observed.

3. Let $f(x) = e^x - 1$

- What is the relative condition number $\kappa(f(x))$? Are there any values of x for which this is ill-conditioned?
- Consider computing $f(x)$ via the following algorithm:

```

1: y = math.e^x
2: return y - 1

```

Is this algorithm stable? Justify your answer

- (c) Let x have the value $9.999999995000000 \times 10^{-10}$, in which case $f(x)$ is equal to 10^{-9} up to 16 decimal places. How many correct digits does the algorithm listed above give you? Is this expected?
- (d) Find a polynomial approximation of $f(x)$ that is accurate to 16 digits for $x = 9.999999995000000 \times 10^{-10}$. Hint: use Taylor series, and remember that 16 digits of accuracy is a relative error, not an absolute one.
- (e) Verify that your answer from part (d) is correct.
- (f) [Optional] How many digits of precision do you have if you do a simpler Taylor series?
- (g) [Fact; no work required] Matlab provides `expm1` and Python provides `numpy.expm1` which are special-purpose algorithms to compute $e^x - 1$ for $x \approx 0$. You could compare your Taylor series approximation with `expm1`.

Soln:

- (a) As we found in class, the condition number is defined as

$$\kappa(f(x)) = \frac{|f'(c)||x|}{|f(x)|}$$

for some c between x and the perturbed value. Since we assume the perturbation is small, we take $c = x$. Thus we get

$$\kappa(f(x)) = \frac{e^x|x|}{|e^x - 1|}.$$

The only point we have to be careful at is $x = 0$ but we can show that $\kappa(f(0)) = 1$. So the condition number is large when x is large.

- (b) For $x \approx 0$, we have $e^x \approx 1$ and we are subtracting nearly equal numbers. This is an unstable algorithm.
- (c) The computation gives us about 8 digits of accuracy. Below, we can see that the computation for $f(x)$ is incorrect in the 9th digit (an 8 instead of a 0). We also see that the relative error is on the order of 10^{-8} which is consistent with 8 digits of accuracy. This is expected as subtracting near equal numbers is unstable.

```
>>> import numpy as np
>>> x = 9.999999995e-10
>>> f = lambda x: np.exp(x) - 1
>>> f(x)
1.000000082740371e-09
>>> fx = 1e-9
>>> ((f(x) - fx))/fx
8.274037093680878e-08
```

- (d) The Taylor remainder term is

$$R_n \leq \frac{1}{(n+1)!} |\xi|^{(n+1)} \leq \frac{1}{(n+1)!} |x|^{(n+1)}$$

for some $\xi \in (-x, x)$. Since we care about at $x = 9.999999995e - 10$, we find that we need two terms in the Taylor expansion in order to achieve our desired error.

- (e,f) The Taylor polynomial $T_2(x) = x(1 + \frac{x}{2})$ satisfies this. Below we see that T_1 gives around 10 digits of accuracy, but T_2 is accurate to machine precision. The value below for x is

```
>>> import numpy as np
>>> x = 9.999999995e-10
>>> T1 = lambda x: x
>>> T2 = lambda x: x*(1 + x/2)
>>> f = 1e-9
>>> (T1(x) - f)/f
>>> (T1(x) - f)/f
-5.000000746056408e-10
>>> (T2(x) - f)/f
0.0
```

4. Practicing using your package

- (a) Create a vector **t** with entries starting at 0 incrementing by $\frac{\pi}{30}$ to π . Then create the vector **y** = cos(**t**).

Write a code that evaluates the following sum:

$$S = \sum_{k=1}^N \mathbf{t}(k) \mathbf{y}(k)$$

Print the statement "the sum is: S with the numerical value of S using the `print` package.

- (b) *Wavy circles*. In one figure plot the parametric curve

$$\begin{aligned} x(\theta) &= R(1 + \delta r \sin(f\theta + p)) \cos(\theta) \\ y(\theta) &= R(1 + \delta r \sin(f\theta + p)) \sin(\theta) \end{aligned}$$

for $0 \leq \theta \leq 2\pi$ and for $R = 1.2$, $\delta r = 0.1$, $f = 15$ and $p = 0$. Make sure to adjust the scale so that the axis have the same scale.

In a second figure use a for loop to plot 10 curves and let with $R = i$, $\delta r = 0.05$, $f = 2 + i$ for the i^{th} curve. Let the value of p be a uniformly distributed random number (look up `random.uniform`) between 0 and 2.

Soln:

- (a) `import numpy as np`
`import math`

```

N = 30
t = np.linspace(0, math.pi, 30)
y = np.cos(t)

S = 0
k = 0
while k < N:
    S = S+ t[k]*y[k]
    k = k +1

print('The sum is', '%16.16e' % S)

```

(b) `import matplotlib.pyplot as plt`
`import numpy as np`
`import math`

```

R = 1.2
delr = 0.1
f = 15.
p = 0.

theta =np.linspace(0, 2*math.pi, 300)

xx = R*(1+delr*np.sin(f*theta + p))*np.cos(theta)
yy = R*(1+delr*np.sin(f*theta + p))*np.sin(theta)

plt.plot(xx,yy)
plt.savefig('prob4bi.pdf')
plt.show()

input()

p = s = np.random.uniform(0,2,1)
delr = 0.05
icount = 1
R = icount
f = 2+icount
xx = R*(1+delr*np.sin(f*theta + p))*np.cos(theta)
yy = R*(1+delr*np.sin(f*theta + p))*np.sin(theta)

plt2 = plt.plot(xx,yy)
while icount <10:
    icount = icount + 1

```

```

R = icount
f = 2+icount
xx = R*(1+delr*np.sin(f*theta + p))*np.cos(theta)
yy = R*(1+delr*np.sin(f*theta + p))*np.sin(theta)
plt2 = plt.plot(xx,yy)

plt2 = plt.savefig('prob4bii.pdf')
plt2 = plt.show()

```

